

Kubernetes Session (19-03-20):

- This is the container Orchestration tool similar to Docker Swarm. Using which we can handle all the production challenges like Load Balancing, Scaling, Rolling updates..etc. Kubernetes was a product of Google and later it was made as open source. The individual machines used in Kubernetes are called as nodes. The Main Node where Kubernetes is installed is called as the Master Node. The Master node is responsible for distributing the workload to slaves. Slave nodes or those nodes which take instructions from the Master and execute. Combination of Master and slaves is Kubernetes Cluster.

KubeADM: This is the application which is installed on the Master and it is responsible for work allocation on the slaves.

KubeAPI: This is an application which is responsible for the slave nodes to remain in contact with the Master.

Kubectl: This application is responsible for executing the Kubernetes Commands.

Kubernetes Objects:

① Pod: This is the smallest object that can be deployed in the Kubernetes cluster. Within a Pod we have a container. Pod acts as layer of encapsulation on a container.

② ReplicaSet: This is a high-level Kubernetes object which can maintain multiple Pod's and collectively perform specific actions on the Pod's like scaling, etc.

- ③ Deployment: This is also a high level Kubernetes object which is used for collective actions on Pod's like Rolling Updates . . . etc
- ④ Service: This works like a router for performing port mapping. It is used for handling the Kubernetes ingress network for External and Internal communication
- ⑤ NameSpace: This is used for creating partitions in the Kubernetes cluster. So that we can deploy multiple applications parallelly.

Kubernetes Setup:

Free setup

- 1 <https://labs.play-with-k8s.com/>
- 2 <https://www.katacoda.com/>

Paid

- 1 sign up for a free Google cloud account
- 2 log in to that account
- 3 click on Navigation Men on top right corner
- 4 click on Kubernetes Engine
- 5 click on Create Cluster
- 6 Go with defaults --> click on Create

Kubernetes Session (20-03-20):

→ Create nginx as a pod on the Kubernetes cluster and name it as webserver check if Pod comes into running condition.

- ① Create a nginx pod

kubectl run --image nginx webserver

- ② To see the list of Pods

kubectl get pods

- ③ To delete this completely

kubectl delete deployment webserver

→ Create a mysql pod in the Kubernetes cluster also pass the necessary environment variables for mysql

- ④ Create a mysql pod

kubectl run --image mysql:5 mydb --env MYSQL_ROOT_PASSWORD=infektaq

- ⑤ To see list of Pods

kubectl get pods

→ To see list of Pods along with the info about the slave where it is running

kubectl get pods -o wide

- ⑥ To delete this completely

kubectl delete deployment mydb

Load Balance:

start httpd with 5 replicas on the Kubernetes cluster

- ① Start httpd with 5 replicas

kubectl run --image httpd appserver --replicas 5

- ② To see the list of Pods

kubectl get pods

- ③ To delete this completely

kubectl delete deployment appserver

Scaling:
Start tomcat with 3 replicas in the kubernetes cluster and later increase it to 7 again scale it down to 2.

① Start tomcat with 3 replicas

Kubectl run --image tomcat mytomcat --replicas 3

② To see ~~if~~ if 3 pods of tomcat are running

Kubectl get pods

③ Increase the replica count to 7

Kubectl scale deployments/mytomcat --replicas 7

④ Check if 7 replicas of tomcat are running

Kubectl get pods

⑤ Decrease the replica count to 2

Kubectl scale deployments/mytomcat --replicas 2

⑥ Delete the tomcat deployment completely

Kubectl delete deployment mytomcat

Definition Files: These definition files used for handling various types of Kubernetes objects. They are created using YAML and they are more reusable. These files have 4 top level fields

① apiVersion:

② kind:

③ Metadata:

④ Spec:

① apiVersion represents the specific version of Kubernetes library that is used for creating a specific object

② kind: This represents kind of Kubernetes object that should be created using this file. Ex: Pod's, service objects .. etc

③ Metadata: Here we store information about the Kubernetes objects like names, labels ..etc. Label section can have any key value pairs

pair
④ Spec: This is where the docker container information is stored like container name, docker image, port mapping etc

kind	ApiVersion
Pod	v1
Service	v1
Name space	v1
Replica set	apps/v1
Replica Deployment	apps/v1

⇒ Create a Pod definition file which will create a Pod named mynginx-pod in which we should have an nginx container with a name mynginx

verm pod-definition-ym)

--
apiVersion: v1

kind: Pod

metadata:

name: nginx-pod

Labels:

author: Intellicq
type: webservice

-Spec

Containers:

- name: mynginx
image: nginx

Aug 11

① To create pods from above file

Kubectl create -f pod-definition.yaml

② To see the list of pods

Kubectl get pods

③ To delete entire deployment

Kubectl delete -f pod-definition.yaml

Kubernetes Session (21-03-20):

→ create a pod definition file which will create ~~a job~~ a MySQL Pod from MySQL container. The name of the Pod should be MySQL:Pod and the labels should be Author=intelleg type=database.

Vim pod-definition2.yaml

apiVersion: v1

Kind: Pod

metadata:

name: mysql-Pod

labels:

author: intelleg

type: database

Spec:

containers:

- name: mydb

image: mysql:5

env:

- name: MYSQL_ROOT_PASSWORD

value: intelleg

...

Wq!

① Kubectl create -f pod-definition2.yaml

② Kubectl get pods

③ Kubectl delete -f pod-definition2.yaml

→ Create a pod-definition file which will start a jenkins container in a Pod called jenkins-pod. Also perform port mapping

Vim pod-definition.yaml

apiVersion: v1

kind: Pod

metadata:

name: jenkins-pod

labels:

author: intelleg

ci: cd

spec:

containers:

- name: jenkins-server

image: jenkins

ports:

- containerPort: 8080

hostPort: 8080

...

Wq!

←
Execution process same as above

ReplicaSet: This is a high level Kubernetes object which is used for maintaining multiple replicas of Pod. The replica count that is given in definition file is treated as desired count and the Replica Set will always try to maintain that desired count. Replica Set uses a field called selector where it will identify elements based on a specific label and it will create the required pods. It also uses a section called template section where we specify actual information about the container.

```
apiVersion  
Kind  
metadata  
spec  
| replicas  
| selector  
| template  
| metadata  
| spec:
```

→ Create a replica set file which will create 3 pods of tomcat
The name of replica set should tomcat-rs labels should be
author: Intelliq type: webserver

```
vim replicas-set.yaml
```

```
---  
apiVersion: apps/v1  
kind: ReplicaSet  
metadata:  
  name: tomcat-rs  
  labels:  
    author: Intelliq  
    type: webserver  
spec:  
  replicas: 3
```

```
  selector:  
    matchLabels:  
      type: webserver
```

```
template  
  template:  
    metadata:  
      name: tomcat-pod  
    labels:  
      type: webserver
```

```
spec:  
  containers:  
    - name: mytomcat  
      image: tomcat  
      ports:  
        - containerPort: 8080  
          hostPort: 9090
```

: WQ!!

Scaling the no. of replicas can be done in 2 ways

- ① By editing the file and editing replace command
- ② By use

① Open the replicas set file

Edit the replicas count to 6

② kubectl replace -f replicas-set.yaml

(or)

Without editing the replicas count in the file

kubectl scale --replicas=2 -f replicas-set.yaml

23-03-20 Class Cancelled

Kubernetes Session (24-03-20):

Deployment: This is also a high level object similar to replicaset. which can be used for managing multiple replicas of a pod. It can also be used for performing scaling activities. And also for performing rolling updates.

→ Create a deployment.yaml file to start 4 replicas of nginx 1.7.9 and later upgrade it to 1.9.1 version

cat deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

labels:

author: Entellegis

type: webserver

spec:

replicas: 4

selector

matchLabels:

type: webserver

```
template:  
  metadata:  
    name: nginx-pod  
    labels:  
      type: web server
```

Spec:

```
  containers:  
    - name: nginx  
      image: nginx:1.7.9  
    ports:  
      - containerPort: 80
```

...

wq!

- ① To create the deployment from the above file
kubectl create -f deployment.yaml
- ② To see if all the 4 pods are running
kubectl get pods
- ③ To see which version of nginx is running
kubectl describe pods pod_name_from_step_2 | less
- ④ To update from nginx:1.7.9 to nginx:1.9.1
kubectl --record deployment.apps/nginx-deployment set image
deployment.v1.apps/nginx-deployment $\begin{cases} \text{deployment name} & \text{nginx=nginx:1.9.1} \\ \text{deployment name} & \text{container name} \end{cases}$
- ⑤ To see the list of pods
kubectl get pods
- ⑥ To see which version of nginx is running
kubectl describe pods pod_name_from_step_5 | less

Kompose: When we implement Docker Compose in Kubernetes it's known as kompose

docker compose + Kubernetes = Kompose

Installing Kompose in Kubernetes Cluster:

① Download kompose

```
CURL -L https://github.com/kubernetes/kompose/releases/download/v1.18.0  
/kompose-linux-amd64 -o kompose
```

② Give execute permissions

```
chmod +x kompose
```

③ Move compose to /usr/local/bin

```
Sudo mv ./kompose /usr/local/bin/kompose
```

④ To check the version of kompose

```
Kompose Version
```

URL: <https://www.digitalocean.com/community/tutorials/how-to-migrate-a-docker-compose-workflow-to-kubernetes>

⇒ Create a docker compose file for starting 1 replica of mysql and 3 replicas of mongoDB

```
Vim docker-compose.yml
```

```
--  
Version: '3'
```

```
services:
```

```
mydb:
```

```
image: mysql:5
```

```
environment
```

```
MYSQL_ROOT_PASSWORD: intellig
```

mywordpress:

image: wordpress

ports:

- 8080: 80

deploy:

replicas: 3

...

wq!

① To start the pods from the above file

Kompose up

② To see the list of pods

kubectl get pods

③ To stop and delete the pods

Kompose down

⇒ Create a dockercompose for starting 1 replica of PostgreSQL
and 4 replicas of nginx

Vim docker-compose.yml

version: '3'

services:

mydb:

image: postgres

environment:

POSTGRES_PASSWORD: Entelliq

webserver:

image: nginx

ports:

- 9090: 80

deploy:

replicas: 4

...

wq!

Follow above 3 steps to run this file

Kubernetes Session (26-03-20)

Service objects in Kubernetes: These Service objects are used when we want elements running in the Kubernetes cluster to communicate to the External network. And also perform network load Balancing. i.e., These Elements should be accessible from any Node, or any slave.

In a Service object we use 3 different ports

- ① TargetPort (ContainerPort)
- ② (Service object Port)
- ③ NodePort (Host Port)

The service objects are also further classified into 3 types

① Cluster IP: This is a default kind of service object that is created. Whenever we create a service definition file and it is used only for internal communication, among the Kubernetes objects.

② Node Port: This type of service object is used when we want to access a Kubernetes object from a specific slave where it is running.

③ LoadBalancer: This service is used when we want to access a Kubernetes object from any of the slaves.

vim pod-definition.yaml (Service is created) (~~not pic taken~~)

vim service1.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    author: intellicq
    type: webserver
```

kubectl create -f pod-definition1.yaml

kubectl create -f service1.yaml

To access nginx from a browser

Give the External IP of slave to 30008

\$ gcloud compute firewall-rules create intellicq --allow tcp:30008

Namespaces in Kubernetes: A Namespace is a partition created in the Kubernetes cluster which is mainly used for separating the services. All objects running in a specific Namespace can communicate with only objects running in the same Namespace.

Vim Namespace.yaml

```
apiVersion: v1
kind: Namespace
metadata:
  name: test-namespace
```

...
wq!

To create Namespace from above file

\$ kubectl create -f namespace.yaml

→ Create a pod-definition file to start a http pod on above created Namespace

Vim pod-definition4.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: httpd-pod
  namespace: test-namespace
  labels:
    - author: entellic
      type: appserver
spec:
  containers:
    - name: myhttpd
      image: httpd
```

...
wq!

To create pod from above file

kubectl create -f pod-definition4.yaml

To see a list of pods running in a specific namespace

kubectl get pods -n test-namespace.

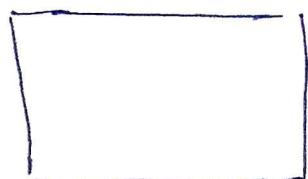
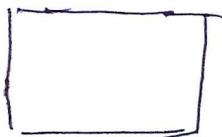
To delete the namespace

kubectl delete namespace test-namespace

Kubernetes Session (27-03-20)!

Kubernetes Project!

This is a voting application which is ~~being~~ created using Python. Here we can cast our vote between two different options. The action that we perform is immediately stored in a temporary database created using Redis. From this Redis database this data is passed to a worker application created using .Net. This worker application processes the data, performs analyses and finally stores it in a persistent database created using PostgreSQL. From here the results can be viewed through an application created using Node.js.



Configuring PyCharm with Kubernetes:

① Download and install Python3

<https://www.python.org/downloads/>

② Download PyCharm

<https://www.jetbrains.com/pycharm/download/#section=windows>

Download the community edition and install it

④ Open PyCharm --> Click on File Menu --> click on setting

Click on Plugins --> Click on setting icon
Click on install from disk --> Browse the path of the plugin that we downloaded

③ Download Kubernetes plugin

<https://plugins.jetbrains.com/plugin/9354-kubernetes-and-openshift-resource-support>

voting-app-pod.yaml

```
-- apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: voting-app-pod
```

```
labels:
```

```
  name: voting-app-pod
```

```
  app: demo-voting-app
```

```
spec:
```

```
  containers:
```

```
    - name: voting-app
```

```
      image: dockersamples/examplevotingapp-vote
```

```
    ports:
```

```
      - containerPort: 80
```

...

redis-pod.yaml

--
apiVersion: v1

kind: Pod

metadata:

name: redis-pod

labels:

name: redis-pod

app: demo-voting-app

spec:

containers:

- name: redis

image: redis

ports:

- ContainerPort: 6379

...

worker-app-pod.yaml

apiVersion: v1

kind: Pod

metadata:

name: worker-app-pod

labels:

name: worker-app-pod

app: demo-voting-app

spec:

containers:

- name: worker

image: docker-samples/example-voting-app-worker

postgres-pod.yaml

apiVersion: v1

kind: Pod

metadata:

name: postgres-pod

labels:

name: postgres-pod

app: demo-voting-app

Spec:

containers:

- name: postgres

image: Postgres:9.4

ports:

- containerPort: 5432

...

Result-app-pod.yaml

apiVersion: v1

kind: Pod

metadata:

name: result-app-pod

labels:

name: result-app-pod

app: demo-voting-app

Spec:

containers:

- name: result-app

image: dockersamples/examplevotingapp-result

ports:

- containerPort: 80

...

Service Definition files

redis-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: redis-service
  labels:
    name: redis-service
    app: demo-voting-app
spec:
  ports:
    - port: 6379
      targetPort: 6379
  selector:
    name: redis-pod
    app: demo-voting-app
  ..
```

postgres-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: postgres-service
  labels:
    name: postgres-service
    app: demo-voting-app
spec:
```

```
  ports:
    - port: 5432
      targetPort: 5432
  selector:
    name: postgres-pod
    app: demo-voting-app
  ..
```

Kubernetes Session (28-03-20)

Voting-app-Service.yaml

```
---  
apiVersion: v1  
kind: Service  
metadata:  
  name: voting-app-service  
  labels:  
    name: voting-app-service  
    app: demo-voting-app  
spec:  
  type: LoadBalancer  
  ports:  
    - port: 80  
      targetPort: 80  
  selector:  
    name: voting-app-pod  
    app: demo-voting-app
```

Result-app-Service.yaml

```
---  
apiVersion: v1  
kind: Service  
metadata:  
  name: result-app-service  
  labels:  
    name: result-app-pod  
    app: demo-voting-app  
spec:
```

```
  type: LoadBalancer  
  ports:  
    - port: 80  
      targetPort: 80  
  selector:  
    name: result-app-pod  
    app: demo-voting-app
```

Docker Swarm

vim voting-app-stack.yml

version: '3'

services:

voting-app:

image: docker.samples/example/voting-app-vote

ports:

- 6767: 80

redis:

image: redis

ports:

- 7777: 6379

worker:

image: docker.samples/example/voting-app-worker

postgres-db:

image: postgres:9.4

environment:

POSTGRES_PASSWORD: inteliq

ports:

- 8888: 5432

result-app:

image: docker.samples/example/voting-app-result

ports:

- 9999: 80

...

docker stack deploy -c voting-app-stack.yml voting-app

Because this entire architecture is running on the production environment & deployed this architecture on Kubernetes cluster. So that we can handle all the challenges like

Load Balancing, High availability ..etc

Since Maintenance of the Kubernetes cluster requires high configuration server, one organisation has decided to migrate this complete environment to Docker Swarm.

I was responsible for the and create the necessary docker compose and docker stack files for setting up this Multi contained architecture in Swarm. Another Major area where we implemented Docker on a large scale is for setting up the testing environment. We have an automation testing team who created test scripts using Selenium. These test scripts are designed to perform cross browser and cross platform testing. Our testing team estimated approximately 800+ Browser and OS combinations to run the test. Since setting up such a huge environment on physical machines or Virtual Machines is very time consuming. Our organisation initially depended on a cloud service provider called ~~sse~~^{sauce} labs but as this was very expensive we did a little bit of POC and identified that we can create the entire ~~the~~ setup using Docker. In just 6 servers we were able to create all these 800+ Browser and OS combinations as docker containers and successfully run the automation test scripts. Now I am actively involved in automating docker related activities using Python.